RAG Configuration

Retrieval Method

vector

Embedding Model

all-mpnet-base-v2

Chunking Strategy

section_based

Prompt Strategy

structured

Top-K Retrieval

5

Gold Standard File

rag_benchmark_answers.json

Defaults = best config from evaluation:
Vector + mpnet + section_based (MRR 1.0)
Structured prompt (ROUGE-L 0.25)

LLM: Qwen/Qwen2.5-0.5B-Instruct

Generation: temp=0.3, top_p=0.9

Mediterranean Cuisine RAG System

Single Query | Benchmark | Evaluation | Exploration Results | About

Evaluation Metrics

Run Evaluation

Retrieval Quality

Precision@5	Recall@5	MRR
0.6533	0.9833	1.0000

Generation Quality

ROUGE-L	BERTScore F1	Faithfulness
0.2522	0.8151	0.6208

Per-Query Breakdown

	Query	ROUGE-L	BERTScore	Faithfulness	P@5
0	Q0	0.1565	0.8138	0.6923	
1	Q1	0.2941	0.8569	0.7	
2	Q2	0.3158	0.8608	0.9231	
3	Q3	0.2082	0.787	0.6154	
4	Q4	0.5116	0.8808	0.8667	
5	Q5	0.324	0.8237	0.8667	
6	Q6	0.3019	0.8491	0.75	
7	Q7	0.2676	0.8306	0.5517	
8	Q8	0.2475	0.7756	0.7727	
9	Q9	0.2482	0.8637	0.4359	

See the **Exploration Results** tab for the full comparison across all 10 retrieval and 3 generation experiments.